

Few-shot object detection Implementation part

Ines BELHAJ MESSAOUD
Jesser BEN HOURIA

National Engineering School of Tunis

19 January 2022



Plan

- 1 Data preparation and visualization
- 2 Implementing YOLOv4
 - Why choosing YOLOv4
- 3 Few-shot implementation
 - MAML

Preparation of our dataset

- ▶ The YOLO implementation is applied on a grapes-dataset that contains images of 5 different classes.
- ▶ The classes represent 5 grape varieties which are the following

Prefix	Variety
CDY	<i>Chardonnay</i>
CFR	<i>Cabernet Franc</i>
CSV	<i>Cabernet Sauvignon</i>
SVB	<i>Sauvignon Blanc</i>
SYH	<i>Syrah</i>

Figure 1: The 5 classes of our dataset

Data preparation



Figure 2: CDY



Figure 3: CSV

Data preparation



Figure 4: CFR



Figure 5: SVB

Data preparation

Initially we had 3 file types in our database :




	CDY_2015.jpg
	CDY_2015.npz
	CDY_2015.txt

Figure 6: File types

- ▶ **jpg** files representing images of grapes.
- ▶ **txt** files representing the bounding boxes coordinates (class - center-x - center-y - width - height).
- ▶ **npz** files representing a format by numpy that provides storage of array data using gzip compression.

Data preparation

We arranged the dataset into two main folders images and labels :

- ▶ Images folder contains the `.jpg` files
- ▶ Labels folder contains the `.txt` files



Figure 7: Data Structure

Data preparation

- It is important to know what each text file represents

1	0	0.1116	0.3026	0.0845	0.2271
2	0	0.5168	0.5084	0.0415	0.0645
3	0	0.6697	0.4912	0.0386	0.0945
4	0	0.0779	0.4835	0.1040	0.1289
5	0	0.0588	0.5813	0.0493	0.0696
6	0	0.0193	0.6216	0.0386	0.1341
7	0	0.1721	0.6066	0.0405	0.0938
8	0	0.1572	0.4187	0.0381	0.0901
9	0	0.1714	0.3502	0.0498	0.0454
10	0	0.2073	0.2916	0.0825	0.0718
11	0	0.4072	0.5806	0.0664	0.1502

Figure 8: Example of a CDY image text file

Data visualization

We have observed the number of images in each category (class) :

```
[ ] n_vimages = {v: len(inst_v) for v, inst_v in instances.items()}  
    n_vimages  
  
{'CDY': 65, 'CFR': 65, 'CSV': 57, 'SVB': 65, 'SYH': 48}
```

Figure 9: Number of images per class

Data visualization




Figure 10: Image distribution per class by percentage

Data visualization

We have also observed the number of bounding boxes in each category (class) :

```
▶ n_vboxes = {v: np.array([n for ii, n in n_iboxes[v].items()]).sum() for v in varietals}  
  n_vboxes
```

```
! {'CDY': 840, 'CFR': 1069, 'CSV': 643, 'SVB': 1316, 'SYH': 563}
```

Figure 11: Number of bounding boxes per class

Data visualization




Figure 12: Boxes distribution per class by percentage

YOLOv4

We chose to implement the [4th YOLO version](#) using both Pytorch and Tensorflow.



Why choosing Pytorch for the training

- ▶ In **PyTorch**, tools and codes are way more imperative and dynamic.
- ▶ The framework is more tightly integrated with Python language.

Why choosing Tensorflow for the visualization

- ▶ [Tensorboard](#) is awesome when it comes to visualization.
- ▶ Useful for debugging and comparison of different training runs.
- ▶ Visualizes the differences between runs

Why choosing YOLOv4




Figure 13: YoloV4 and YoloV3

Results of y

- ▶ YOLOv4 is an improvement on the YOLOv3 algorithm. The mean average precision (mAP) improved by 10% and the number of frames per second improved by 12%.
- ▶ The YOLOv4 architecture has 4 different blocks : The backbone, the neck, the dense prediction, and the sparse prediction.

Evaluation Metrics

- ▶ We focused on evaluating our yolov4 model with 3 main metrics :
- ▶ mAP : AP (Average precision) is a popular metric in measuring the accuracy of object detectors like Faster R-CNN, SSD, etc. Average precision computes the average precision value for recall value over 0 to 1.
- ▶ Precision
- ▶ recall

YoloV3 results

- ▶ With YoloV3, with 300 epochs, we have obtained a **precision** = 0.67 and a **recall** = 0.47.
- ▶ The average **IoU** is equal to 47.8.
- ▶ With 300 epochs, we have obtained a $\text{mAp} = 59.39 \%$

Results

- ▶ All the training results and the test results are downloaded in the inference folder.
- ▶ The tensorboard also provides all the train and test metrics obtained with every Epoch.

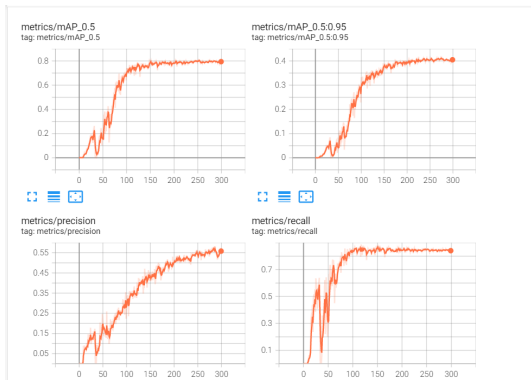


Figure 14: Result example with 300 epochs

Results

YOLO loss function is broken into three parts:

- ▶ The one responsible for finding the bounding-box coordinates (cls-loss : cross entropy loss)
- ▶ The bounding-box score prediction (giou-loss)
- ▶ The class-score prediction (obj-loss) : Mean squared error loss

Train Results

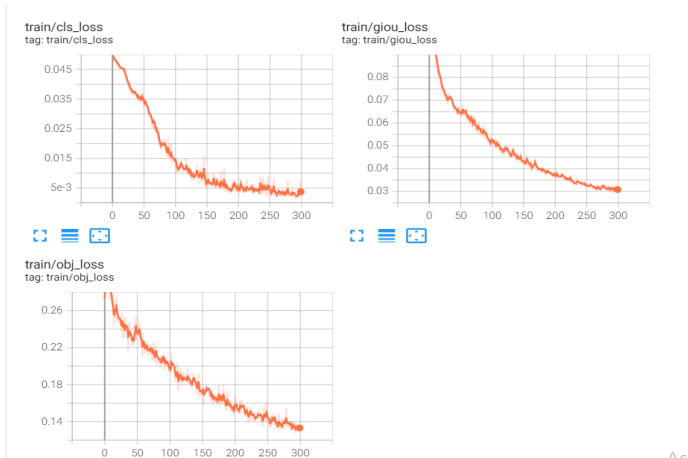


Figure 15: Train Result example with 300 epochs

Test Results

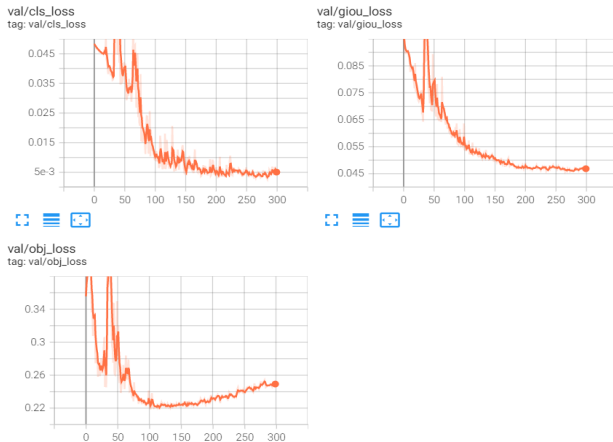


Figure 16: Test Result example with 300 epochs

Interpretation

- ▶ With 100 epochs we have noticed an underfitting : the map of the training set was remarkably low
- ▶ When we increased the number of epochs to 300 we obtained a map of 0.8.

FEW-SHOT implementation

First step of Few-shot implementation

- For the Few-shot training part, we chose to implement the MAML algorithm.
- MAML's main goal is optimizing the model parameters so that a small number of gradient steps would produce a maximum effective behaviour on a new task.




Figure 1. Diagram of our model-agnostic meta-learning algorithm (MAML), which optimizes for a representation θ that can quickly adapt to new tasks.

First step of Few-shot implementation

- ▶ The **meta-optimization across tasks** is performed with stochastic gradient descent.
- ▶ The **model parameters** are noted as θ and β is the meta step size

Algorithm 2 MAML for Few-Shot Supervised Learning

Require: $p(\mathcal{T})$: distribution over tasks

Require: α, β : step size hyperparameters

```
1: randomly initialize  $\theta$ 
2: while not done do
3:   Sample batch of tasks  $\mathcal{T}_i \sim p(\mathcal{T})$ 
4:   for all  $\mathcal{T}_i$  do
5:     Sample  $K$  datapoints  $\mathcal{D} = \{\mathbf{x}^{(j)}, \mathbf{y}^{(j)}\}$  from  $\mathcal{T}_i$ 
6:     Evaluate  $\nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$  using  $\mathcal{D}$  and  $\mathcal{L}_{\mathcal{T}_i}$  in Equation (2)
       or (3)
7:     Compute adapted parameters with gradient descent:
        $\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$ 
8:     Sample datapoints  $\mathcal{D}'_i = \{\mathbf{x}^{(j)}, \mathbf{y}^{(j)}\}$  from  $\mathcal{T}_i$  for the
       meta-update
9:   end for
10:  Update  $\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$  using each  $\mathcal{D}'_i$ 
    and  $\mathcal{L}_{\mathcal{T}_i}$  in Equation 2 or 3
11: end while
```

Figure 17: Model-Agnostic Meta-Learning

First step of Few-shot implementation

Algorithm 2 MAML for Few-Shot Supervised Learning

Require: $p(\mathcal{T})$: distribution over tasks

Require: α, β : step size hyperparameters

```

1: randomly initialize  $\theta$ 
2: while not done do
3:   Sample batch of tasks  $\mathcal{T}_i \sim p(\mathcal{T})$ 
4:   for all  $\mathcal{T}_i$  do
5:     Sample  $K$  datapoints  $\mathcal{D} = \{\mathbf{x}^{(j)}, \mathbf{y}^{(j)}\}$  from  $\mathcal{T}_i$ 
6:     Evaluate  $\nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$  using  $\mathcal{D}$  and  $\mathcal{L}_{\mathcal{T}_i}$  in Equation (2)
       or (3)
7:     Compute adapted parameters with gradient descent:
        $\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$ 
8:     Sample datapoints  $\mathcal{D}'_i = \{\mathbf{x}^{(j)}, \mathbf{y}^{(j)}\}$  from  $\mathcal{T}_i$  for the
       meta-update
9:   end for
10:  Update  $\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$  using each  $\mathcal{D}'_i$ 
    and  $\mathcal{L}_{\mathcal{T}_i}$  in Equation 2 or 3
11: end while
```

Figure 18: MAML algorithm

```

#In our case the model is a yolov4 model
model=ConvolutionalNeuralNetwork(out_features=7) #we propose a 7-way setting
meta_optimizer=torch.optim.Adam(model.parameters(), lr=1e-3)
#Sample batch of tasks/ Line_4
for task in batch:

    #Ligne_5 and Line_6
    train_inputs, train_targets=task['support']
    test_inputs, test_targets=task['query']

    train_logits=model(train_input) #Evaluate the model
    inner_loss=F.cross_entropy(train_logits, train_target)
    model.zero_grad()

    #Line_7
    #estimate the gradients on the inner loss on the model.meta_params
    grads=torch.autograd.grad(inner, model.meta_params(), create_graph=True)
    params=OrderedDict() #define a dictionary of params

    #iterate on our model.meta_params and we translate the equation
    #(do a gradient step with a time step size)
    for (name, param), grad in zip (model.meta_named_pars(), grads):
        params[name]=param - step_size * grad

    #Line_8 to line_10
    test_logits=model(test_input, params=params)

    #the sum of all the losses over all the tasks
    outer_loss+=F.cross_entropy(test_logits, test_target)
    outer_loss.backward() #apply the backward pass over all the tasks
    meta_optimizer.step() #apoly the optimizer with the specific step size
```

Figure 19: Pytorch implementation of MAML algorithm

THANK YOU FOR YOUR ATTENTION